

Parallel Minimum degree ordering

When factoring sparse matrices, new non-zero entries are created. This phenomenon, known as “fill-in”, has a significant impact on performance. This amount of extra non-zero entries created depends on column ordering. Determining good column orderings is therefore an extremely precomputing phase of any sparse matrix factorization.

1 Elimination game

The **elimination game** consists in simulating the factorization itself while monitoring newly created non-zero entries. Based on this knowledge, it is possible to design heuristics intending to reduce the amount of fill-in.

Sparse matrix row-column dependencies are represented using a graph structure $\mathcal{G} = (V, E)$. Each node $v \in V$ corresponds to a column in the sparse matrix. Every edge $(u, v) \in E$ between two nodes u and v represents a dependency between these two columns, i.e. a non zero entry in (u, v) and (v, u) entries of the sparse matrix.

At every step, the elimination game processes as follows:

1. Pick a node v and eliminate it.
2. Create a clique between all nodes adjacent to v

2 Minimum degree algorithm

2.1 Parallelization schemes

- algorithms not explicitly changing the graph
- algorithms changing the graph

3 Finding indistinguishable nodes when computing reachable set

4 Performance

Input: v is the node eliminated at step s . R_v is its reachable set.
 $marker$ and $label$ arrays, tag

```

 $label(v) = s$ 
 $indistCount = 0$ 
 $tag = tag + 1$ 
 $tag_v = tag$ 
forall the node  $u \in R_v$  do
    |  $mask(u) = tag_v$ 
end
forall the node  $t \in R_v$  do
    |  $tag = tag + 1$ 
    |  $indist, deg(t) \leftarrow$ 
    |  $update\_degree(t, v, deg(v), label, marker, tag, tag_v, mask)$ 
    | if  $indist$  then
    | |  $s = s + 1$ 
    | |  $label(t) = s$ 
    | |  $indistCount \leftarrow indistCount + 1$ 
    | end
end
forall the node  $t \in R_v$  do
    | if  $label(t) = 0$  then
    | |  $deg(t) \leftarrow deg(t) - indistCount$ 
    | end
end

```

Algorithm 1: Sketch of the MDO algorithm calling $update_degree$

Input:

1. u , node of which we're computing the reachable set (starting point of the exploration).
2. v , node eliminated at current step.
3. $deg(v)$, degree of v .
4. $label$, array of size n indicating if a node has been labeled or not.
5. $marker$, array of size n used to mark explored nodes with value tag .
6. tag_v , special tag value used to mark nodes in R_v .
7. $mask$, array of size n used to mark nodes in R_v with tag_v .

Output:

1. $indist$, boolean indicating if v and u are indistinguishable.
2. $\bar{deg}(u)$, updated degree of u after the elimination of v .

```

 $\bar{deg}(u) \leftarrow deg(v) - 1$ 
 $explore \leftarrow \{u\}$ 
 $indist \leftarrow true$ 
 $count \leftarrow 1$ 
forall the node  $t$  in  $explore$  do
    forall the node  $x$  in  $Adj_t$  do
        if  $marker(x) \neq tag$  then
            if  $label(x) \neq 0$  then
                if  $x \neq v$  then
                     $explore \leftarrow explore \cup \{x\}$ 
                end
            end
        else
            if  $mask(x) \neq tag_v$  then
                 $indist \leftarrow false$ 
                 $\bar{deg}(u) \leftarrow \bar{deg}(u) + 1$ 
            end
            else
                 $count \leftarrow count + 1$ 
            end
        end
         $marker(x) = tag$ 
    end
end
end
if  $indist = true$  AND  $count + 1 \neq deg(v)$  then
     $indist \leftarrow false$ 
end

```

Algorithm 2: *update_degree*

P	Reach	Reach & comp.	Reach, comp. & mass.	Explicit	Explicit & mass.
1		4.9444	5.2108	3.5210	0.2097
4		1.9232	2.1678	1.4694	0.2074
8		1.1575	1.4402	0.9910	0.1702
16		0.8842	1.3877	0.8215	0.1937
24		0.8289	1.3861	0.8287	0.2133